



# CS 4240: Compilers

Lecture 18: Register Allocation Review

Instructor: Vivek Sarkar


([vsarkar@gatech.edu](mailto:vsarkar@gatech.edu))

March 25, 2019

# ANNOUNCEMENTS & REMINDERS

---

- » **Project 2 due by 11:59pm on Wednesday, April 3rd**
  - » 15% of course grade
  - » Next lecture (March 27) will be a help session for Project 2
- » **FINAL EXAM: Wednesday, May 1, 2:40 pm - 5:30 pm**
  - » 30% of course grade





# Register allocation (Lecture 12)

- Assume a 3-address code intermediate representation with an unbounded number of virtual/symbolic registers
- For each virtual register  $\mathbf{r}$ , want to choose a physical machine register to store  $\mathbf{r}$ 's values
- Intermediate step: determine the *live ranges* of each virtual register  $\mathbf{r}$

# A Simple Example

```
s1 ← 2  
s2 ← 4  
s3 ← s1 + s2  
s4 ← s3 + 1  
s5 ← s1 * s2  
s6 ← s4 * 2
```



- **Live**: variable will be used before it is overwritten
- **Dead**: variable will be overwritten before it is used

# A Simple Example (contd)

Intermediate  
Code w/ Symbolic  
Registers

```
s1 ← 2  
s2 ← 4  
s3 ← s1 + s2  
s4 ← s3 + 1  
s5 ← s1 * s2  
s6 ← s4 * 2
```

Machine  
Code w/ Physical  
Registers

```
r1 ← 2  
r2 ← 4  
r3 ← r1 + r2  
r3 ← r3 + 1  
r1 ← r1 * r2  
r2 ← r3 * 2
```

# Building the Interference Graph

---

What is an “interference” ? (or conflict)

- » Two values **interfere** if there exists an operation where both are simultaneously live
- » If  $x$  and  $y$  interfere, they cannot occupy the same register

To compute interferences, we must know where values are “live”

The interference graph,  $G_I = (N_I, E_I)$

- » Nodes in  $G_I$  represent values, or live ranges
- » Edges in  $G_I$  represent individual interferences
  - For  $x, y \in N_I$ ,  $\langle x, y \rangle \in E_I$  iff  $x$  and  $y$  interfere
- » A  $k$ -coloring of  $G_I$  can be mapped into an allocation to  $k$  registers


# Graph Coloring

---


## The problem

A graph  $G$  is said to be  $k$ -colorable iff the nodes can be labeled with integers  $1 \dots k$  so that no edge in  $G$  connects two nodes with the same label

## Examples



2-colorable



3-colorable

Each color can be mapped to a distinct physical register

## Computing LIVE Sets (live ranges)

---

A value  $v$  is live at program point  $p$  if  $\exists$  a path from  $p$  to some use of  $v$  along which  $v$  is not re-defined

Data-flow problems are expressed as simultaneous equations

$$\text{LIVEOUT}(b) = \bigcup_{s \in \text{succ}(b)} \text{LIVEIN}(s)$$

$$\text{LIVEIN}(b) = \text{UEVAR}(b) \cup (\text{LIVEOUT}(b) - \text{VARKILL}(b))$$

where

$\text{UEVAR}(b)$  is the set of names used in block  $b$  before being defined in  $b$  (Upwards Exposed Variables)

$\text{VARKILL}(b)$  is the set of variables assigned in  $b$

Solve the equations using a fixed-point iterative scheme



# Computing LIVE Sets

---

The compiler can solve these equations with an iterative algorithm

```
WorkList  $\leftarrow$  { all blocks }  
while ( WorkList  $\neq \emptyset$  )  
    remove a block b from WorkList  
    Compute LIVEOUT(b)  
    Compute LIVEIN(b)  
    if LIVEIN(b) changed  
        then add pred (b) to WorkList
```

The Worklist Iterative  
Algorithm

Why does this work?

- LIVEOUT, LIVEIN  $\subseteq 2^{\text{Names}}$
- UEVAR, VARKILL are constants for b
- Equations are monotone
- Finite # of additions to sets  
 $\Rightarrow$  will reach a fixed point !

Speed of convergence depends on the order in which blocks are “removed” & their sets recomputed

The world's quickest introduction to data-flow analysis !

# Register Allocation

1. Determine live ranges for each symbolic register
2. Determine overlapping ranges (**interference**)
3. Compute the benefit of keeping each live range in a register (**spill cost**)
4. Try to assign each live range to a machine register (**allocation**).  
If needed, **spill** or **split** live range
5. Generate code, including spills

# **Extended Linear Scan: an Alternate Foundation for Global Register Allocation**

**Compiler Construction (CC) conference, 2007**

Vivek Sarkar (IBM T.J. Watson Research Center)

Rajkishore Barik (IBM India Research Laboratory)

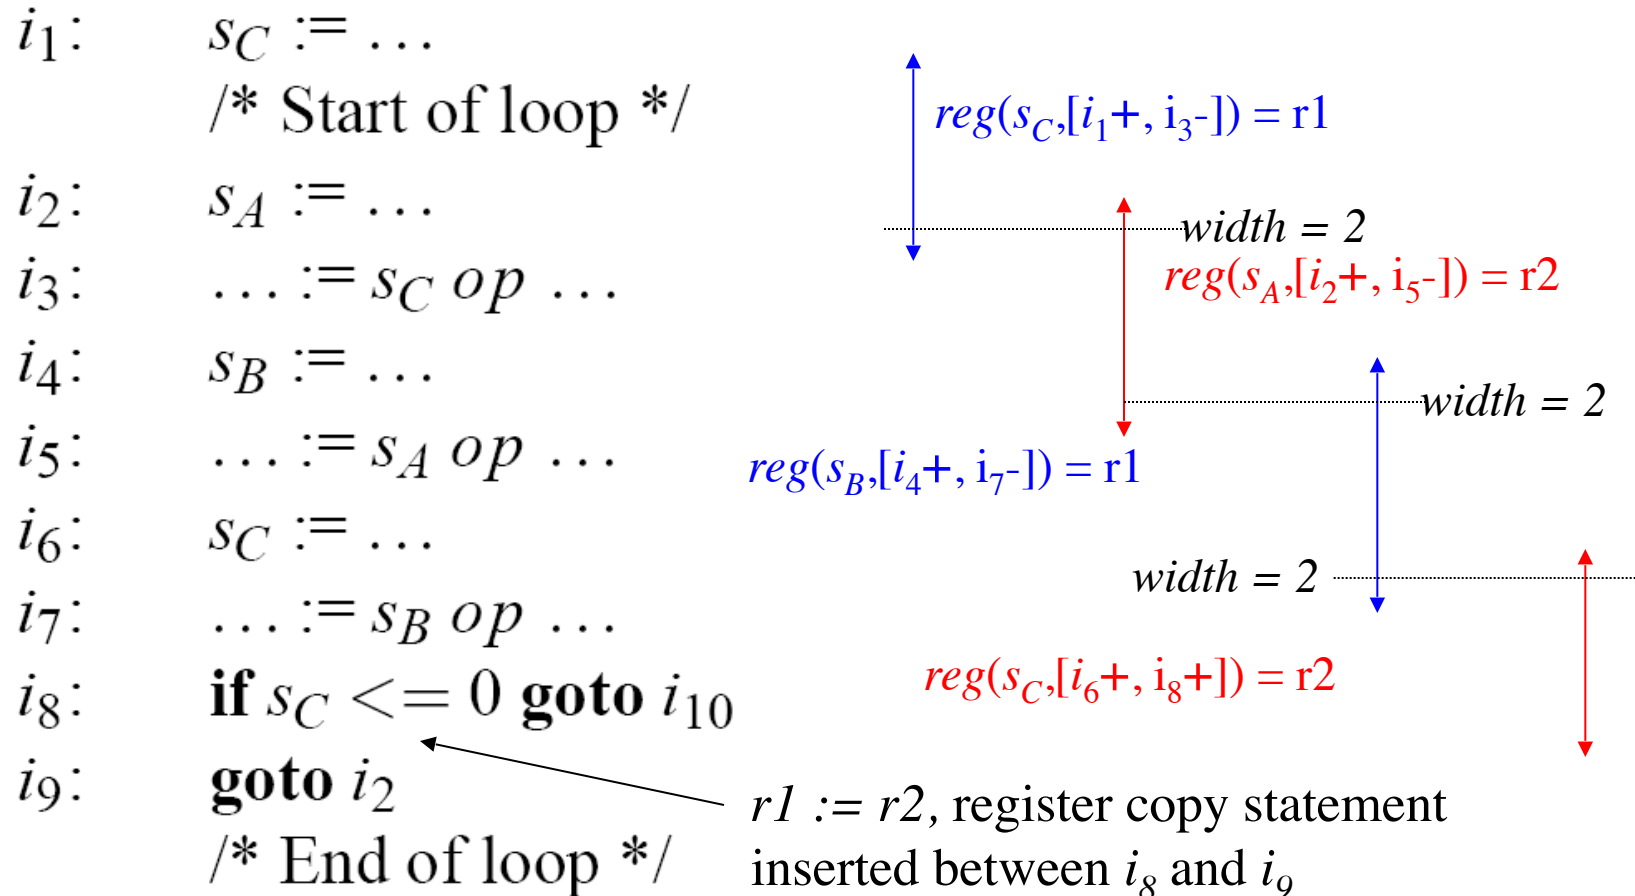
# Summary

- Register allocation continues in importance since first FORTRAN compiler
- Motivation: theoretical and practical limitations of Graph Coloring
- Two variants of Register Allocation problem addressed in this paper:
  - Spill Free Register Allocation (SFRA) – decision problem
  - Register Allocation with Total Spills (RATS) – optimization problem
- Approach includes possibility of inserting register moves to change register assignments at different program points (regMoves=true/false)
- Two new Extended Linear Scan (ELS) register allocation algorithms
  - $ELS_0$  for SFRA problem (only for regMoves=true)
  - $ELS_1$  for RATS problem (two versions based on regMoves=true/false)
- Comparison of  $ELS_1$  (regMoves=true) with Graph Coloring for RATS problem on eight SPECint2000 benchmarks
  - 15x – 68x speedups in *compile-time* relative to Graph Coloring (compile-time measured on largest functions in the eight benchmarks)
  - Up to 5.8% speedup (average = 2.3%) in *run-time* relative to GC (run-time measured for entire benchmark)
  - No coalescing or live range splitting heuristics implemented for this study

# Spill Free Register Allocation


- Given a set of symbolic registers,  $\mathbf{R}$ , and  $k$  physical registers, determine if it is possible to assign each symbolic register  $s$  to a physical register,  $reg(s, P)$  at each program point  $P$  where  $s$  is live.
  - If so, report the register assignments, including any register-to-register copy statements that need to be inserted.
  - If not, report that no feasible solution exists.
- Two program points are defined for each instruction  $i$ 
  - $i-$  = point at which input operands are read
  - $i+$  = point at which output operands are written

# Example of a Spill Free Register Allocation with Register Copies ( $k=2$ )



# Limitations of Graph Coloring Approach

- Graph Coloring is a more limited problem than Spill Free Register Allocation
  - Finding a  $k$ -coloring of an Interference Graph will lead to a feasible solution to the SFRA problem, but the converse is not true
- Non-linear space complexity for Interference Graph
  - $O(R^2)$  worst case,  $O(R^{1.5})$  observed in practice
- Graph coloring is NP-hard (without a constant performance bound), but SFRA with `regMoves=true` can be solved optimally in linear time
- Interference Graph for previous example



# Overview of $ELS_0$ Extended Linear Scan Algorithm for SFRA

- **Steps performed**
  - Compute  $\text{count}[P] = \#$  intervals live at each program point  $P$  that is an interval end point
  - If any  $\text{count}[P]$  is  $> k$  then *return* false
  - Assign registers to intervals
  - Insert register moves as needed on each control flow edge
  - *Return* true
- **Complexity:** linear in size of IR and number of intervals
- **Correctness:**  $ELS_0$  computes an exact solution to the SFRA decision problem



# Space Savings for $ELS_0$ relative to Graph Coloring

Function	$ S $	$ IG $	$ \mathcal{I} $	SCF
164.gzip.build_tree	161	2301	261	11.3%
175.vpr.try_route	254	2380	445	18.7%
181.mcf.sort_basket	138	949	226	22.7%
186.crafty.InputMove	122	1004	219	21.8%
197.parser.list_links	352	9090	414	4.5%
254.gap.SyFgets	547	7661	922	12.0%
256.bzip2.sendMTFValues	256	2426	430	17.7%
300.twolf.closepins	227	5105	503	9.8%

**Table 1.** Compile-time overheads for functions with the largest interference graphs in SPECint2000 benchmarks.  $|S|$  = # symbolic registers,  $|IG|$  = # nodes and edges in Interference Graph,  $|\mathcal{I}|$  = # intervals in interval set, Space Compression Factor (SCF) =  $|\mathcal{I}|/|IG|$ .

# Register Allocation with Total Spills

- Given a set of symbolic registers,  $\mathbf{R}$ ,  $k$  physical registers, and estimated execution frequency  $freq[P]$  for each program point  $P$ , compute the following:
  1.  $spilled(s)$ , which indicates if  $s$  is to be spilled, and
  2. for each symbolic register  $s$  with  $spilled(s) = false$ , a register assignment,  $reg(s, P)$  at each point  $P$  where  $s$  is live.
- Two versions of the RATS problem:
  - $regMoves = false$ . Optimization goal is to minimize frequency-weighted *spill cost*.
  - $regMoves = true$ . Optimization goal is to minimize combined frequency-weighted *spill cost* and *register moves* (using architecture-specific relative weightages)
- RATS optimization problem is NP-hard for both  $regMoves = false$  and *true* cases

# Overview of ELS<sub>1</sub> Heuristic for RATS

- Steps performed for regMoves = true case:
  - Compute count[P] = # intervals live at each program point P that is an interval end point
  - Select point P with count[P] > k in decreasing order of freq[P]
    - Push symbolic register live at P with smallest spill cost on stack
    - Repeat until count[P] ≤ k for all program points
  - Resurrection: “unspill” symbolic registers on stack that no longer need to be spilled
  - Register assignment as in ELS<sub>0</sub>
    - Assign registers to remaining intervals (guaranteed to be spill free)
    - Insert register moves as needed on each control flow edge
- Modification for regMoves = false
  - “All or nothing” heuristic for register assignment avoids copies
  - Additional spills may need to be inserted
- Complexity: linear in size of IR and number of intervals

# Summary

- Register allocation continues in importance since first FORTRAN compiler
- Motivation: theoretical and practical limitations of Graph Coloring
- Two variants of Register Allocation problem addressed in this paper:
  - Spill Free Register Allocation (SFRA) – decision problem
  - Register Allocation with Total Spills (RATS) – optimization problem
- Approach includes possibility of inserting register moves to change register assignments at different program points (regMoves=true/false)
- Two new Extended Linear Scan (ELS) register allocation algorithms
  - $ELS_0$  for SFRA problem (only for regMoves=true)
  - $ELS_1$  for RATS problem (two versions based on regMoves=true/false)
- Comparison of  $ELS_1$  (regMoves=true) with Graph Coloring for RATS problem on eight SPECint2000 benchmarks
  - 15x – 68x speedups in *compile-time* relative to Graph Coloring (compile-time measured on largest functions in the eight benchmarks)
  - Up to 5.8% speedup (average = 2.3%) in *run-time* relative to GC (run-time measured for entire benchmark)
  - No coalescing or live range splitting heuristics implemented for this study